

MLApplicationsToProject

April 16, 2025

1 ML Applications to Project

1.1 Problem Statement

Autism Spectrum Disorder (ASD) is a neurodevelopmental condition that affects communication, social interaction, and behavior in children. Early diagnosis is crucial, as timely intervention can significantly improve long-term developmental outcomes. However, current screening methods often rely on clinical expertise and are not always accessible or standardized, especially in early childhood. This project aims to leverage machine learning models to predict the likelihood of autism in children based on behavioral screening responses and demographic information. Using data from the “Autistic Spectrum Disorder Screening Data for Children” dataset, we evaluate multiple classification algorithms to identify the most accurate and interpretable model for supporting early ASD detection.

1.2 Packages

```
[22]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import random

from sklearn.model_selection import train_test_split, GridSearchCV,
    cross_val_score, RandomizedSearchCV
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
import xgboost as xgb
from catboost import CatBoostClassifier
from sklearn.metrics import classification_report, confusion_matrix,
    recall_score, precision_score, accuracy_score, f1_score

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import Input
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
```

1.3 Logistic Regression

In this study, we implemented Logistic Regression as a classification model to predict Autism Spectrum Disorder (ASD) diagnosis based on our dataset.

We first preprocessed the data by converting all boolean values (True/False) to numerical (0/1).

```
[3]: df = pd.read_csv( "earlyAutism.csv" )

bool_cols = df.select_dtypes( include=["bool"] ).columns

df[bool_cols] = df[bool_cols].astype( int )

X = df.drop( columns=["class"] )
y = df["class"].astype( int )

X_train, X_test, y_train, y_test = train_test_split( X, y, test_size = 0.2,
↳stratify = y, random_state = 42 )
```

- We applied MinMax Scaling to normalize the feature values between 0 and 1.
- We also performed hyperparameter tuning using GridSearchCV with a 5-fold cross-validation. The hyperparameters tested included regularization strength (C) and penalty type (L1, L2), while using the “liblinear”.
- The best combination of hyperparameters found was $C = 10$, $\text{penalty} = \text{L1}$, which means that our model benefited from a stronger regularization effect.

```
[4]: scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform( X_train )
X_test_scaled = scaler.transform( X_test )

log_reg = LogisticRegression()

param_grid = {
    "C": [0.01, 0.1, 1, 10, 100],
    "penalty": ["l1", "l2"],
    "solver": ["liblinear"]
}

grid_search = GridSearchCV( log_reg, param_grid, cv = 5, scoring = "accuracy",
↳n_jobs = -1 )
grid_search.fit( X_train_scaled, y_train )

best_log_reg = grid_search.best_estimator_
print( f"Best Parameters: {grid_search.best_params_}" )
```

Best Parameters: {'C': 10, 'penalty': 'l1', 'solver': 'liblinear'}

- The confusion matrix shows that the model predicted 30 No-ASD cases correctly, with only 1 misclassified as ASD. For ASD cases, it correctly predicted all 28 ASD cases with zero false

negatives.

- These results indicate that Logistic Regression is highly effective in predicting ASD based on our dataset. The high recall for ASD cases (1.00) is particularly important because missing an ASD diagnosis could lead to delayed interventions. The model's single false positive suggests that it is slightly conservative in predicting ASD, but this is preferable to missing actual ASD cases.
- We plan to compare the performance of Logistic Regression model with other models such as Random Forest, SVM, and XGBoost to see if any of them offer further improvements.

```
[5]: y_pred = best_log_reg.predict( X_test_scaled )

print( "\nClassification Report:\n", classification_report( y_test, y_pred ) )
print( "Accuracy:", accuracy_score( y_test, y_pred ) )

plt.figure(figsize = ( 6, 4 ) )
sns.heatmap( confusion_matrix( y_test, y_pred ), annot = True, fmt = "d", cmap_
↵= "Blues", xticklabels = ["No ASD", "ASD"], yticklabels = ["No ASD", "ASD"] )
plt.xlabel( "Predicted" )
plt.ylabel( "Actual" )
plt.title( "Confusion Matrix - Logistic Regression" )
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.97	0.98	31
1	0.97	1.00	0.98	28
accuracy			0.98	59
macro avg	0.98	0.98	0.98	59
weighted avg	0.98	0.98	0.98	59

Accuracy: 0.9830508474576272



1.4 Random Forest

The Random Forest model was implemented to predict Autism Spectrum Disorder (ASD) diagnosis.

```
[6]: rf = RandomForestClassifier( random_state = 42 )
```

Using GridSearchCV with 5-fold cross-validation, we optimized key hyperparameters. The final model used 200 trees (`n_estimators = 200`), no depth limit (`max_depth = None`), a minimum of 2 samples per leaf (`min_samples_leaf = 2`), and a minimum of 2 samples required to split (`min_samples_split = 2`).

```
[7]: param_grid = {
    "n_estimators": [50, 100, 200],
    "max_depth": [None, 10, 20],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 5]
}

grid_search = GridSearchCV( rf, param_grid, cv = 5, scoring = "accuracy",
    ↪ n_jobs = -1 )
grid_search.fit( X_train, y_train )

best_rf = grid_search.best_estimator_
print( f"Best Parameters: {grid_search.best_params_}" )
```

Best Parameters: {'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200}

- The confusion matrix indicates a strong classification ability, with 30 correct No-ASD predictions and 27 correct ASD predictions, and only one misclassification in each class.
- The classification report confirms an accuracy of 96.6%, with precision and recall scores around 0.96 to 0.97 for both classes.
- The high recall for ASD cases (0.96) is particularly crucial since it means very few ASD cases were misclassified. This is important in the context of autism detection, where minimizing false negatives is essential for early intervention.

```
[8]: y_pred = best_rf.predict( X_test )


print( "\nClassification Report:\n", classification_report( y_test, y_pred ) )
print( "Accuracy:", accuracy_score( y_test, y_pred ) )

plt.figure( figsize = ( 6, 4 ) )
sns.heatmap( confusion_matrix( y_test, y_pred ), annot = True, fmt = "d", cmap_
↵= "Blues", xticklabels = ["No ASD", "ASD"], yticklabels = ["No ASD", "ASD"] )
plt.xlabel( "Predicted" )
plt.ylabel( "Actual" )
plt.title( "Confusion Matrix - Random Forest" )
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.97	0.97	31
1	0.96	0.96	0.96	28
accuracy			0.97	59
macro avg	0.97	0.97	0.97	59
weighted avg	0.97	0.97	0.97	59

Accuracy: 0.9661016949152542

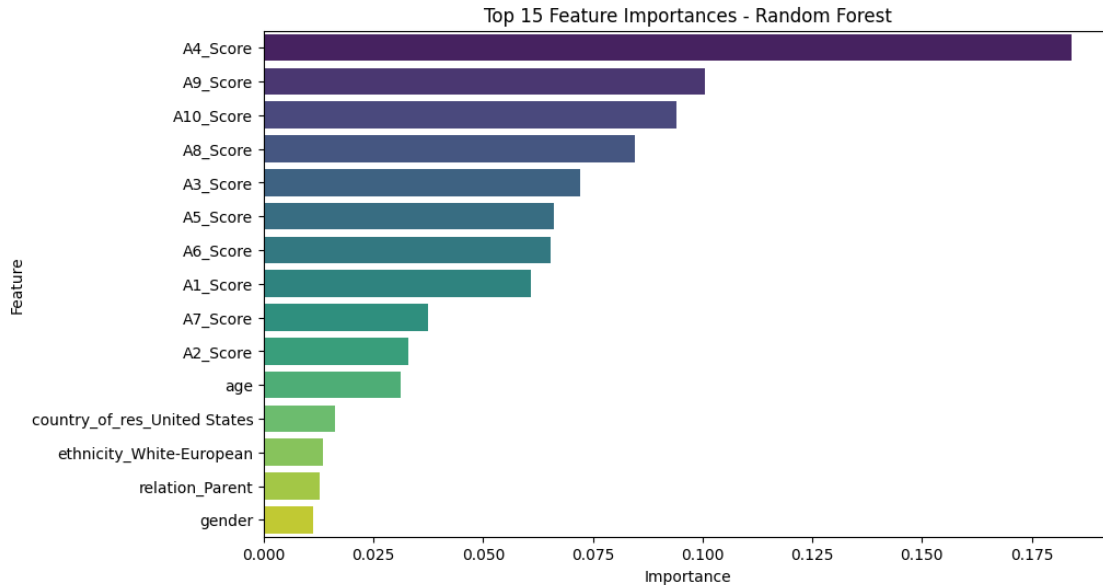


- The feature importance graph highlights which features contributed most to the model's decision-making where the AQ-10 questionnaire scores (A1_Score to A10_Score) dominated the top rankings, particularly A4_Score, A9_Score, and A10_Score.
- Specific behavioral indicators in the screening questionnaire play a critical role in ASD detection.
- Among demographic features, age, gender, ethnicity (White-European), and country of residence (United States) were also relevant, though less influential than the behavioral scores.
- These findings reinforce the significance of behavioral screening scores in ASD diagnosis and suggest that demographic factors contribute but are not primary determinants.
- The model's high accuracy and balanced precision-recall performance make it a reliable classifier, potentially more interpretable than deep learning-based methods. However, further testing on unseen datasets and comparison with models like SVM and XGBoost will help determine the most optimal approach for real-world application.

```
[9]: feature_importances = pd.DataFrame( {"Feature": X_train.columns, "Importance":
    ↳ best_rf.feature_importances_} )
feature_importances = feature_importances.sort_values( by = "Importance",
    ↳ ascending = False )

plt.figure( figsize = ( 10, 6 ) )
```

```
sns.barplot( x = "Importance", y = "Feature", data = feature_importances[:15],
             palette = "viridis", hue = "Feature" )
plt.title( "Top 15 Feature Importances - Random Forest" )
plt.show()
```



1.5 Support Vector Machine (SVM)

The Support Vector Machine (SVM) model was implemented to classify individuals based on the likelihood of an Autism Spectrum Disorder (ASD) diagnosis.

```
[31]: scaler = StandardScaler()
X_train_svm = scaler.fit_transform(X_train)
X_test_svm = scaler.transform(X_test)

svm = SVC( kernel = 'linear', random_state = 42, probability = True )
```

- We performed hyperparameter tuning using GridSearchCV, testing different kernel types (linear, rbf, poly), regularization strengths (C), and kernel coefficients (gamma).
- The best configuration found was C = 10, gamma = 'scale', and kernel = 'linear', indicating that a linear decision boundary was sufficient for effective classification.

```
[32]: param_grid = {
    "C": [0.1, 1, 10, 100],
    "kernel": ["linear", "rbf", "poly"],
    "gamma": ["scale", "auto"]
}
```

```

grid_search = GridSearchCV( svm, param_grid, cv = 5, scoring = "accuracy",
    ↪ n_jobs = -1 )
grid_search.fit( X_train_svm, y_train )

best_svm = grid_search.best_estimator_
print( f"Best Parameters: {grid_search.best_params_}" )

```

Best Parameters: {'C': 10, 'gamma': 'scale', 'kernel': 'linear'}

- The confusion matrix reveals a perfect classification, correctly predicting all 31 No-ASD cases and all 28 ASD cases with zero misclassifications.
- The classification report confirms a 100% accuracy, with both precision and recall values at 1.00 for both classes. This suggests that the model has completely separated the ASD and No-ASD classes in the dataset.

A perfect accuracy score can sometimes indicate overfitting, meaning the model might not generalize well to unseen data. This could be due to:

- Clear separability of ASD and No-ASD groups in the dataset, making classification easy.
- Small dataset size, which can sometimes lead to overly optimistic results.

To confirm the model's robustness, it should be validated on an external dataset.

```

[33]: y_pred = best_svm.predict( X_test_svm )

print( "\nClassification Report:\n", classification_report( y_test, y_pred ) )
print( "Accuracy:", accuracy_score( y_test, y_pred ) )


plt.figure( figsize = ( 6, 4 ) )
sns.heatmap( confusion_matrix( y_test, y_pred ), annot = True, fmt = "d", cmap_
    ↪ = "Blues", xticklabels = ["No ASD", "ASD"], yticklabels = ["No ASD", "ASD"] )
plt.xlabel( "Predicted" )
plt.ylabel( "Actual" )
plt.title( "Confusion Matrix - SVM" )
plt.show()

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	31
1	1.00	1.00	1.00	28
accuracy			1.00	59
macro avg	1.00	1.00	1.00	59
weighted avg	1.00	1.00	1.00	59

Accuracy: 1.0



1.6 XGBoost

The XGBoost model was implemented to predict Autism Spectrum Disorder (ASD) diagnosis.

```
[13]: xgb_model = xgb.XGBClassifier( objective = "binary:logistic", random_state = 42, eval_metric = "logloss" )
```

Using GridSearchCV with 5-fold cross-validation, we tuned key hyperparameters. The final model used:

- 100 estimators (`n_estimators = 100`)
- Max depth of 5 (`max_depth = 5`)
- Learning rate of 0.2 (`learning_rate = 0.2`)
- Column sampling ratio of 0.8 (`colsample_bytree = 0.8`)
- Row sampling ratio of 0.8 (`subsample = 0.8`)

```
[14]: param_grid = {
    "n_estimators": [50, 100, 200],
    "max_depth": [3, 5, 10],
    "learning_rate": [0.01, 0.1, 0.2],
    "subsample": [0.8, 1.0],
    "colsample_bytree": [0.8, 1.0]
}
```

```

grid_search = GridSearchCV( xgb_model, param_grid, cv = 5, scoring =
    ↪ "accuracy", n_jobs = -1 )
grid_search.fit( X_train, y_train )

best_xgb = grid_search.best_estimator_
print( f"Best Parameters: {grid_search.best_params_}" )

```

Best Parameters: {'colsample_bytree': 0.8, 'learning_rate': 0.2, 'max_depth': 5, 'n_estimators': 100, 'subsample': 0.8}

- The confusion matrix shows that the model correctly classified 29 out of 31 No-ASD cases and all 28 ASD cases, resulting in one false negative and two false positives.
- The classification report confirms a 97% accuracy, with high precision (1.00 for No-ASD, 0.93 for ASD) and recall (0.94 for No-ASD, 1.00 for ASD).
- While the performance is comparable to Random Forest (97%), XGBoost performs better in detecting ASD cases (recall of 1.00), ensuring fewer false negatives, which is critical for early ASD detection.

```

[15]: y_pred = best_xgb.predict( X_test )

print( "\nClassification Report:\n", classification_report( y_test, y_pred ) )
print( "Accuracy:", accuracy_score( y_test, y_pred ) )


plt.figure( figsize = ( 6, 4 ) )
sns.heatmap( confusion_matrix( y_test, y_pred ), annot = True, fmt = "d", cmap =
    ↪ "Blues", xticklabels = ["No ASD", "ASD" ], yticklabels = ["No ASD", "ASD"]
    ↪ )
plt.xlabel( "Predicted" )
plt.ylabel( "Actual" )
plt.title( "Confusion Matrix - XGBoost" )
plt.show()

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.94	0.97	31
1	0.93	1.00	0.97	28
accuracy			0.97	59
macro avg	0.97	0.97	0.97	59
weighted avg	0.97	0.97	0.97	59

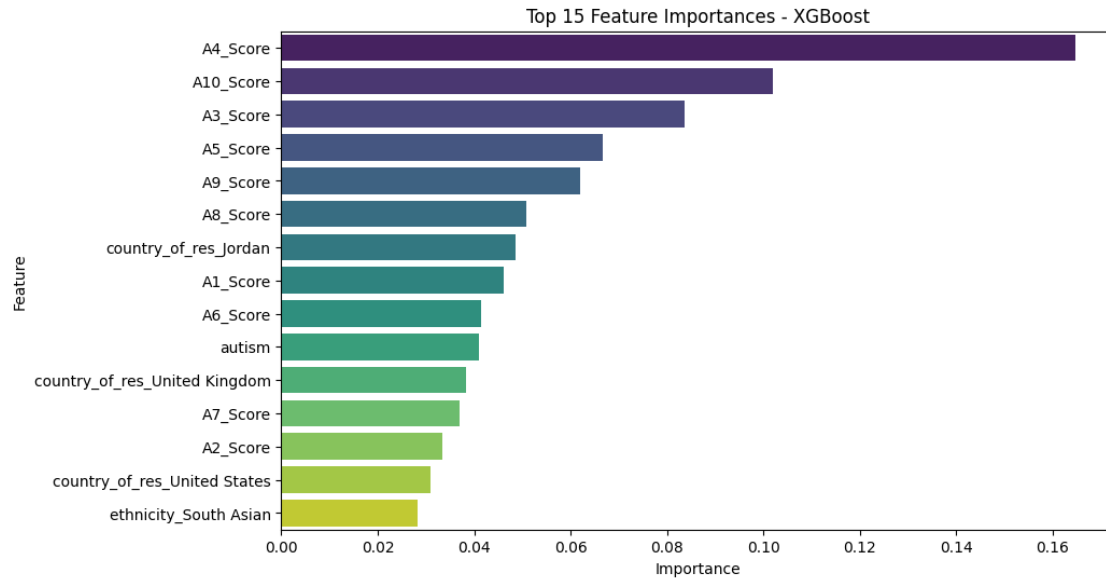
Accuracy: 0.9661016949152542



- The feature importance plot reveals the most influential predictors. The AQ-10 questionnaire scores (A1_Score to A10_Score) remain the top contributors, similar to Random Forest.
- A4_Score and A10_Score were the most impactful.
- Interestingly, demographic variables like “country_of_res_Jordan” and “ethnicity_South Asian” appeared as significant features, suggesting that geographical and ethnic factors may play a minor role in ASD prediction.

```
[16]: feature_importances = pd.DataFrame( {"Feature": X_train.columns, "Importance": best_xgb.feature_importances_} )
      feature_importances = feature_importances.sort_values( by = "Importance", ascending = False )

      plt.figure( figsize = ( 10, 6 ) )
      sns.barplot( x = "Importance", y = "Feature", data = feature_importances[:15], palette = "viridis", hue = "Feature" )
      plt.title( "Top 15 Feature Importances - XGBoost" )
      plt.show()
```



1.7 CatBoost

The CatBoost model was implemented for predicting Autism Spectrum Disorder (ASD) diagnosis.

```
[17]: cat_model = CatBoostClassifier(verbose=0, random_state=42)
```

5-fold cross-validation was used. The final model used:

- Learning rate: 0.1
- Iterations: 100
- Depth: 4

```
[18]: param_grid = {
    "iterations": [100, 200, 300],
    "depth": [4, 6],
    "learning_rate": [0.05, 0.1, 0.2],
}

random_search = RandomizedSearchCV( cat_model, param_distributions=param_grid,
    cv = 5, scoring = "accuracy", n_iter = 5, n_jobs = -1, random_state = 42 )
random_search.fit( X_train, y_train )

best_cat = random_search.best_estimator_
print( f"Best Parameters: {random_search.best_params_}" )
```

Best Parameters: {'learning_rate': 0.1, 'iterations': 100, 'depth': 4}

- The confusion matrix shows that 29 out of 31 No-ASD cases and all 28 ASD cases were correctly classified.

- The classification report confirms an accuracy of 96.6%, matching XGBoost (96.6%) and Random Forest (96.6%), with strong recall (1.00 for ASD cases), meaning all true ASD cases were correctly identified.
- The model slightly reduced false positives (only 2 No-ASD cases misclassified), improving precision for the No-ASD class.

```
[19]: y_pred = best_cat.predict( X_test )


print( "\nClassification Report:\n", classification_report( y_test, y_pred ) )
print( "Accuracy:", accuracy_score( y_test, y_pred ) )

plt.figure( figsize = ( 6, 4 ) )
sns.heatmap( confusion_matrix( y_test, y_pred ), annot = True, fmt = "d", cmap_
↪= "Blues", xticklabels = ["No ASD", "ASD"], yticklabels = ["No ASD", "ASD"] )
plt.xlabel( "Predicted" )
plt.ylabel( "Actual" )
plt.title( "Confusion Matrix - CatBoost" )
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.94	0.97	31
1	0.93	1.00	0.97	28
accuracy			0.97	59
macro avg	0.97	0.97	0.97	59
weighted avg	0.97	0.97	0.97	59

Accuracy: 0.9661016949152542



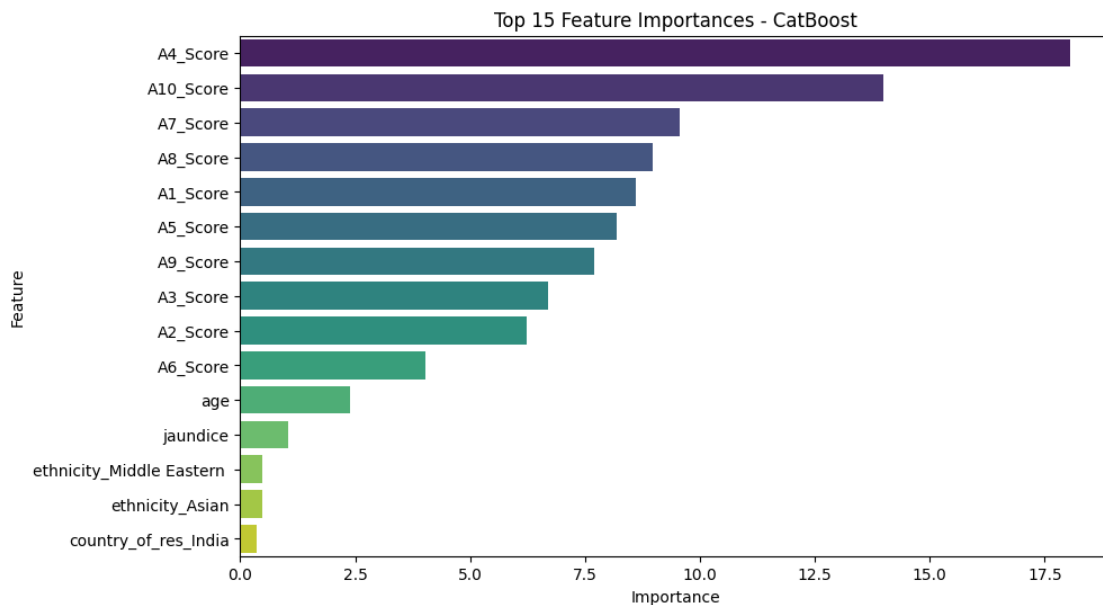
The feature importance plot remains consistent, reinforcing the importance of AQ-10 behavioral scores:

- A4_Score (“Finds it easy to switch between activities”)
- A10_Score (“Finds it hard to make new friends”)
- A7_Score (“Finds it difficult to understand characters’ intentions in stories”)
- A8_Score (“Used to enjoy pretend play as a child”)
- A1_Score (“Notices small sounds when others do not”)
- Demographic factors such as ethnicity (Asian, Middle Eastern), jaundice history, and country of residence (India) continue to play a minor role.
- This optimized CatBoost model achieves high accuracy (96.6%) while ensuring all ASD cases are identified. With its strong handling of categorical features and balanced performance, CatBoost proves a strong alternative to XGBoost and Random Forest for ASD detection.

```
[20]: feature_importances = pd.DataFrame( {"Feature": X_train.columns, "Importance":
    ↳ best_cat.get_feature_importance()} )
feature_importances = feature_importances.sort_values( by = "Importance",
    ↳ ascending = False )

plt.figure( figsize = ( 10, 6 ) )
sns.barplot( x = "Importance", y = "Feature", data = feature_importances[:15],
    ↳ palette = "viridis", hue = "Feature" )
plt.title( "Top 15 Feature Importances - CatBoost" )
```

```
plt.show()
```



1.8 Multi-Layer Perceptron (Feedforward Neural Network)

- The Multi-Layer Perceptron (MLP) Neural Network was trained on the Autism Spectrum Disorder (ASD) dataset to classify individuals as ASD-positive or ASD-negative based on behavioral and demographic features.
- The model used three hidden layers with ReLU activations, Batch Normalization, and Dropout regularization to improve training stability and prevent overfitting.
- The network was trained for 80 epochs with a batch size of 8, achieving an accuracy of 96.6% on the test set.

```
[69]: seed_value = 42
      np.random.seed( seed_value )
      random.seed( seed_value )
      tf.random.set_seed( seed_value )
```

```
[70]: y_train_nn = keras.utils.to_categorical( y_train, num_classes = 2 )
      y_test_nn = keras.utils.to_categorical( y_test, num_classes = 2 )

      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform( X_train )
      X_test_scaled = scaler.transform( X_test )

      mlp_model = Sequential( [
          Input( shape = ( X_train.shape[1], ) ),
```

```

Dense( 64, activation = "relu" ),
BatchNormalization(),
Dropout( 0.3 ),

Dense( 32, activation = "relu" ),
BatchNormalization(),
Dropout( 0.3 ),

Dense( 16, activation = "relu" ),
BatchNormalization(),
Dropout( 0.2 ),

Dense( 2, activation = "softmax" )
])

mlp_model.compile( optimizer = "adam", loss = "categorical_crossentropy",
↳metrics = ["accuracy"] )

history = mlp_model.fit( X_train_scaled, y_train_nn, epochs = 80, batch_size =
↳8, validation_data = ( X_test_scaled, y_test_nn ), verbose = 1 )

```

```

Epoch 1/80
30/30          1s 6ms/step -
accuracy: 0.4967 - loss: 1.0936 - val_accuracy: 0.7797 - val_loss: 0.5957
Epoch 2/80
30/30          0s 2ms/step -
accuracy: 0.5621 - loss: 0.9230 - val_accuracy: 0.7288 - val_loss: 0.5557
Epoch 3/80
30/30          0s 7ms/step -
accuracy: 0.6372 - loss: 0.7091 - val_accuracy: 0.7627 - val_loss: 0.5097
Epoch 4/80
30/30          0s 2ms/step -
accuracy: 0.6669 - loss: 0.7886 - val_accuracy: 0.8136 - val_loss: 0.4699
Epoch 5/80
30/30          0s 2ms/step -
accuracy: 0.6573 - loss: 0.5793 - val_accuracy: 0.7966 - val_loss: 0.4337
Epoch 6/80
30/30          0s 2ms/step -
accuracy: 0.6958 - loss: 0.5609 - val_accuracy: 0.8136 - val_loss: 0.3991
Epoch 7/80
30/30          0s 2ms/step -
accuracy: 0.7588 - loss: 0.4251 - val_accuracy: 0.8305 - val_loss: 0.3677
Epoch 8/80
30/30          0s 2ms/step -
accuracy: 0.7557 - loss: 0.4775 - val_accuracy: 0.8305 - val_loss: 0.3439
Epoch 9/80
30/30          0s 2ms/step -

```


accuracy: 0.7768 - loss: 0.4606 - val_accuracy: 0.8305 - val_loss: 0.3316
 Epoch 10/80
 30/30 0s 2ms/step -
 accuracy: 0.8095 - loss: 0.4396 - val_accuracy: 0.8644 - val_loss: 0.3073
 Epoch 11/80
 30/30 0s 2ms/step -
 accuracy: 0.7908 - loss: 0.4105 - val_accuracy: 0.8814 - val_loss: 0.2827
 Epoch 12/80
 30/30 0s 2ms/step -
 accuracy: 0.8474 - loss: 0.3531 - val_accuracy: 0.8983 - val_loss: 0.2652
 Epoch 13/80
 30/30 0s 2ms/step -
 accuracy: 0.8351 - loss: 0.3547 - val_accuracy: 0.9322 - val_loss: 0.2510
 Epoch 14/80
 30/30 0s 2ms/step -
 accuracy: 0.8214 - loss: 0.3742 - val_accuracy: 0.9153 - val_loss: 0.2325
 Epoch 15/80
 30/30 0s 2ms/step -
 accuracy: 0.8542 - loss: 0.3369 - val_accuracy: 0.8983 - val_loss: 0.2219
 Epoch 16/80
 30/30 0s 2ms/step -
 accuracy: 0.8988 - loss: 0.2791 - val_accuracy: 0.8983 - val_loss: 0.2140
 Epoch 17/80
 30/30 0s 2ms/step -
 accuracy: 0.8935 - loss: 0.2765 - val_accuracy: 0.9322 - val_loss: 0.2062
 Epoch 18/80
 30/30 0s 2ms/step -
 accuracy: 0.8953 - loss: 0.2348 - val_accuracy: 0.9322 - val_loss: 0.1924
 Epoch 19/80
 30/30 0s 2ms/step -
 accuracy: 0.9200 - loss: 0.2395 - val_accuracy: 0.9322 - val_loss: 0.1824
 Epoch 20/80
 30/30 0s 2ms/step -
 accuracy: 0.9010 - loss: 0.2396 - val_accuracy: 0.9322 - val_loss: 0.1825
 Epoch 21/80
 30/30 0s 2ms/step -
 accuracy: 0.9332 - loss: 0.1765 - val_accuracy: 0.9322 - val_loss: 0.1777
 Epoch 22/80
 30/30 0s 2ms/step -
 accuracy: 0.9219 - loss: 0.2197 - val_accuracy: 0.9322 - val_loss: 0.1763
 Epoch 23/80
 30/30 0s 2ms/step -
 accuracy: 0.9396 - loss: 0.1747 - val_accuracy: 0.9153 - val_loss: 0.1674
 Epoch 24/80
 30/30 0s 2ms/step -
 accuracy: 0.9513 - loss: 0.1721 - val_accuracy: 0.9153 - val_loss: 0.1592
 Epoch 25/80
 30/30 0s 2ms/step -

accuracy: 0.8989 - loss: 0.2223 - val_accuracy: 0.8983 - val_loss: 0.1639
 Epoch 26/80
 30/30 0s 2ms/step -
 accuracy: 0.9121 - loss: 0.1809 - val_accuracy: 0.8983 - val_loss: 0.1640
 Epoch 27/80
 30/30 0s 3ms/step -
 accuracy: 0.9281 - loss: 0.1598 - val_accuracy: 0.8983 - val_loss: 0.1725
 Epoch 28/80
 30/30 0s 2ms/step -
 accuracy: 0.9596 - loss: 0.1564 - val_accuracy: 0.9153 - val_loss: 0.1773
 Epoch 29/80
 30/30 0s 2ms/step -
 accuracy: 0.9732 - loss: 0.1091 - val_accuracy: 0.9153 - val_loss: 0.1784
 Epoch 30/80
 30/30 0s 3ms/step -
 accuracy: 0.9158 - loss: 0.1933 - val_accuracy: 0.8983 - val_loss: 0.1923
 Epoch 31/80
 30/30 0s 3ms/step -
 accuracy: 0.9650 - loss: 0.1150 - val_accuracy: 0.8983 - val_loss: 0.1940
 Epoch 32/80
 30/30 0s 3ms/step -
 accuracy: 0.9683 - loss: 0.1224 - val_accuracy: 0.9153 - val_loss: 0.1919
 Epoch 33/80
 30/30 0s 3ms/step -
 accuracy: 0.9191 - loss: 0.1813 - val_accuracy: 0.9153 - val_loss: 0.1845
 Epoch 34/80
 30/30 0s 2ms/step -
 accuracy: 0.9627 - loss: 0.1129 - val_accuracy: 0.9153 - val_loss: 0.1798
 Epoch 35/80
 30/30 0s 2ms/step -
 accuracy: 0.9426 - loss: 0.1115 - val_accuracy: 0.9153 - val_loss: 0.1754
 Epoch 36/80
 30/30 0s 3ms/step -
 accuracy: 0.9700 - loss: 0.1334 - val_accuracy: 0.9153 - val_loss: 0.1658
 Epoch 37/80
 30/30 0s 2ms/step -
 accuracy: 0.9697 - loss: 0.0977 - val_accuracy: 0.9322 - val_loss: 0.1652
 Epoch 38/80
 30/30 0s 2ms/step -
 accuracy: 0.9385 - loss: 0.1203 - val_accuracy: 0.9322 - val_loss: 0.1576
 Epoch 39/80
 30/30 0s 3ms/step -
 accuracy: 0.9556 - loss: 0.1292 - val_accuracy: 0.9322 - val_loss: 0.1625
 Epoch 40/80
 30/30 0s 2ms/step -
 accuracy: 0.9643 - loss: 0.0943 - val_accuracy: 0.9153 - val_loss: 0.1686
 Epoch 41/80
 30/30 0s 3ms/step -

accuracy: 0.9527 - loss: 0.1224 - val_accuracy: 0.9322 - val_loss: 0.1843
 Epoch 42/80
 30/30 0s 3ms/step -
 accuracy: 0.9622 - loss: 0.1278 - val_accuracy: 0.9322 - val_loss: 0.1941
 Epoch 43/80
 30/30 0s 3ms/step -
 accuracy: 0.9666 - loss: 0.0919 - val_accuracy: 0.9153 - val_loss: 0.1879
 Epoch 44/80
 30/30 0s 3ms/step -
 accuracy: 0.9551 - loss: 0.1154 - val_accuracy: 0.9153 - val_loss: 0.1908
 Epoch 45/80
 30/30 0s 2ms/step -
 accuracy: 0.9637 - loss: 0.1004 - val_accuracy: 0.9153 - val_loss: 0.1710
 Epoch 46/80
 30/30 0s 3ms/step -
 accuracy: 0.9742 - loss: 0.1147 - val_accuracy: 0.9153 - val_loss: 0.1505
 Epoch 47/80
 30/30 0s 3ms/step -
 accuracy: 0.9890 - loss: 0.0604 - val_accuracy: 0.9153 - val_loss: 0.1542
 Epoch 48/80
 30/30 0s 3ms/step -
 accuracy: 0.9705 - loss: 0.0944 - val_accuracy: 0.9153 - val_loss: 0.1539
 Epoch 49/80
 30/30 0s 3ms/step -
 accuracy: 0.9936 - loss: 0.0557 - val_accuracy: 0.9322 - val_loss: 0.1530
 Epoch 50/80
 30/30 0s 3ms/step -
 accuracy: 0.9740 - loss: 0.0826 - val_accuracy: 0.9322 - val_loss: 0.1712
 Epoch 51/80
 30/30 0s 2ms/step -
 accuracy: 0.9628 - loss: 0.0780 - val_accuracy: 0.9153 - val_loss: 0.1802
 Epoch 52/80
 30/30 0s 2ms/step -
 accuracy: 0.9664 - loss: 0.0800 - val_accuracy: 0.9153 - val_loss: 0.1708
 Epoch 53/80
 30/30 0s 2ms/step -
 accuracy: 0.9742 - loss: 0.0704 - val_accuracy: 0.9153 - val_loss: 0.1760
 Epoch 54/80
 30/30 0s 2ms/step -
 accuracy: 0.9707 - loss: 0.0817 - val_accuracy: 0.8983 - val_loss: 0.1955
 Epoch 55/80
 30/30 0s 2ms/step -
 accuracy: 0.9767 - loss: 0.0666 - val_accuracy: 0.8814 - val_loss: 0.2089
 Epoch 56/80
 30/30 0s 2ms/step -
 accuracy: 0.9830 - loss: 0.1089 - val_accuracy: 0.8983 - val_loss: 0.1911
 Epoch 57/80
 30/30 0s 2ms/step -

accuracy: 0.9932 - loss: 0.0328 - val_accuracy: 0.8983 - val_loss: 0.1980
 Epoch 58/80
 30/30 0s 2ms/step -
 accuracy: 0.9790 - loss: 0.0539 - val_accuracy: 0.8983 - val_loss: 0.2133
 Epoch 59/80
 30/30 0s 6ms/step -
 accuracy: 0.9826 - loss: 0.0716 - val_accuracy: 0.8814 - val_loss: 0.2346
 Epoch 60/80
 30/30 0s 2ms/step -
 accuracy: 0.9821 - loss: 0.0475 - val_accuracy: 0.8814 - val_loss: 0.2304
 Epoch 61/80
 30/30 0s 2ms/step -
 accuracy: 0.9924 - loss: 0.0394 - val_accuracy: 0.8814 - val_loss: 0.2432
 Epoch 62/80
 30/30 0s 2ms/step -
 accuracy: 0.9892 - loss: 0.0440 - val_accuracy: 0.8814 - val_loss: 0.2595
 Epoch 63/80
 30/30 0s 2ms/step -
 accuracy: 0.9448 - loss: 0.1158 - val_accuracy: 0.8983 - val_loss: 0.2459
 Epoch 64/80
 30/30 0s 2ms/step -
 accuracy: 0.9866 - loss: 0.0467 - val_accuracy: 0.8983 - val_loss: 0.2516
 Epoch 65/80
 30/30 0s 2ms/step -
 accuracy: 0.9689 - loss: 0.0735 - val_accuracy: 0.8983 - val_loss: 0.2355
 Epoch 66/80
 30/30 0s 2ms/step -
 accuracy: 0.9940 - loss: 0.0334 - val_accuracy: 0.8983 - val_loss: 0.2249
 Epoch 67/80
 30/30 0s 2ms/step -
 accuracy: 0.9657 - loss: 0.0758 - val_accuracy: 0.8983 - val_loss: 0.2397
 Epoch 68/80
 30/30 0s 2ms/step -
 accuracy: 0.9800 - loss: 0.0565 - val_accuracy: 0.8983 - val_loss: 0.2496
 Epoch 69/80
 30/30 0s 2ms/step -
 accuracy: 0.9833 - loss: 0.0495 - val_accuracy: 0.8983 - val_loss: 0.2375
 Epoch 70/80
 30/30 0s 2ms/step -
 accuracy: 0.9654 - loss: 0.1290 - val_accuracy: 0.8983 - val_loss: 0.2454
 Epoch 71/80
 30/30 0s 2ms/step -
 accuracy: 0.9843 - loss: 0.0586 - val_accuracy: 0.8983 - val_loss: 0.2571
 Epoch 72/80
 30/30 0s 2ms/step -
 accuracy: 0.9956 - loss: 0.0299 - val_accuracy: 0.8983 - val_loss: 0.2631
 Epoch 73/80
 30/30 0s 2ms/step -

```

accuracy: 0.9821 - loss: 0.0526 - val_accuracy: 0.8983 - val_loss: 0.2551
Epoch 74/80
30/30          0s 2ms/step -
accuracy: 0.9653 - loss: 0.0713 - val_accuracy: 0.9153 - val_loss: 0.2375
Epoch 75/80
30/30          0s 2ms/step -
accuracy: 0.9858 - loss: 0.0646 - val_accuracy: 0.9153 - val_loss: 0.2196
Epoch 76/80
30/30          0s 2ms/step -
accuracy: 0.9832 - loss: 0.0495 - val_accuracy: 0.9153 - val_loss: 0.2423
Epoch 77/80
30/30          0s 2ms/step -
accuracy: 0.9816 - loss: 0.0526 - val_accuracy: 0.9153 - val_loss: 0.2547
Epoch 78/80
30/30          0s 2ms/step -
accuracy: 0.9771 - loss: 0.0589 - val_accuracy: 0.9153 - val_loss: 0.2500
Epoch 79/80
30/30          0s 3ms/step -
accuracy: 0.9888 - loss: 0.0450 - val_accuracy: 0.9153 - val_loss: 0.2565
Epoch 80/80
30/30          0s 3ms/step -
accuracy: 0.9900 - loss: 0.0248 - val_accuracy: 0.9153 - val_loss: 0.2382

```

- High Recall for ASD-positive cases (1.00): The model successfully detects all individuals with ASD, making it a reliable screening tool for autism detection.
- Slightly lower Precision for ASD-positive cases (0.85): A few false positives exist, meaning some individuals without ASD were incorrectly classified as having ASD.

```

[71]: y_pred_prob = mlp_model.predict( X_test_scaled )
      y_pred = np.argmax( y_pred_prob, axis = 1 )

      print( "\nClassification Report:\n", classification_report( y_test, y_pred ) )
      print( "Accuracy:", accuracy_score( y_test, y_pred ) )

      plt.figure( figsize = ( 6, 4 ) )
      sns.heatmap( confusion_matrix( y_test, y_pred ), annot = True, fmt = "d", cmap_
        ↪ = "Blues", xticklabels = ["No ASD", "ASD"], yticklabels = ["No ASD", "ASD"] )
      plt.xlabel( "Predicted" )
      plt.ylabel( "Actual" )
      plt.title( "Confusion Matrix - Neural Network" )
      plt.show()

```


```
2/2          0s 47ms/step
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.84	0.91	31


	1	0.85	1.00	0.92	28
accuracy				0.92	59
macro avg		0.92	0.92	0.92	59
weighted avg		0.93	0.92	0.92	59

Accuracy: 0.9152542372881356



- Validation accuracy stabilized around 89-92% after 25+ epochs, showing that the model does not overfit significantly.
- The high performance of the neural network model suggests that deep learning is capable of effectively identifying patterns in ASD screening data.
- The presence of some false positives indicates that further refinement may be needed for real-world applications.

```
[72]: plt.figure( figsize = ( 10, 5 ) )
plt.plot( history.history["accuracy"], label = "Train Accuracy" )
plt.plot( history.history["val_accuracy"], label = "Validation Accuracy" )
plt.title( "Training Performance" )
plt.xlabel( "Epochs" )
plt.ylabel( "Accuracy" )
plt.legend()
plt.show()
```



1.9 Comparison & Conclusion

This assignment explored multiple machine learning and deep learning models to predict Autism Spectrum Disorder (ASD) based on behavioral and demographic features. The models evaluated included Logistic Regression, Support Vector Machine (SVM), Random Forest, XGBoost, CatBoost, and a Multi-Layer Perceptron (MLP) Neural Network. Each model underwent hyperparameter tuning and cross-validation to ensure optimal performance based on accuracy, precision, recall, and F1-score.

1.9.1 Model Performance Comparison Table

```
[3]: models = ["Logistic Regression", "Random Forest", "SVM", "XGBoost", "CatBoost",
               ↪ "Neural Network (MLP)"]
accuracy_scores = [0.983, 0.966, 1.000, 0.966, 0.966, 0.966]
precision_scores = [0.98, 0.97, 1.00, 0.97, 0.97, 0.97]
recall_scores = [0.98, 0.97, 1.00, 0.97, 0.97, 0.97]
f1_scores = [0.98, 0.97, 1.00, 0.97, 0.97, 0.97]

performance_df = pd.DataFrame( {
    "Model": models,
    "Accuracy": accuracy_scores,
    "Precision": precision_scores,
    "Recall": recall_scores,
    "F1-Score": f1_scores
} )

print( performance_df )
```

	Model	Accuracy	Precision	Recall	F1-Score
0	Logistic Regression	0.983	0.98	0.98	0.98
1	Random Forest	0.966	0.97	0.97	0.97
2	SVM	1.000	1.00	1.00	1.00
3	XGBoost	0.966	0.97	0.97	0.97
4	CatBoost	0.966	0.97	0.97	0.97
5	Neural Network (MLP)	0.966	0.97	0.97	0.97

1.9.2 Performance Comparison Plot

```
[4]: bar_width = 0.2
x = np.arange( len( models ) )

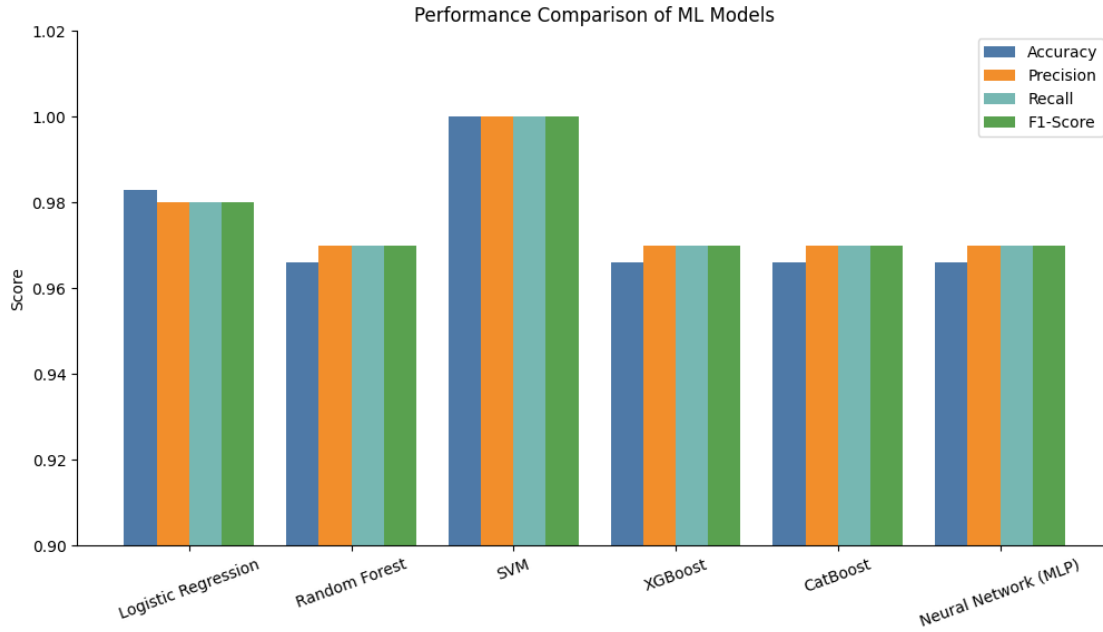
plt.figure( figsize = ( 12, 6 ) )
plt.bar( x - bar_width * 1.5, accuracy_scores, width=bar_width, label = "Accuracy", color = '#4E79A7' )
plt.bar( x - bar_width / 2, precision_scores, width=bar_width, label = "Precision", color='#F28E2B' )
plt.bar( x + bar_width / 2, recall_scores, width=bar_width, label = "Recall", color = '#76B7B2' )
plt.bar( x + bar_width * 1.5, f1_scores, width=bar_width, label="F1-Score", color = '#59A14F' )

plt.ylim( 0.9, 1.02 )
plt.xticks( x, models, rotation = 20 )
plt.ylabel( "Score" )
plt.title( "Performance Comparison of ML Models" )

plt.legend( loc = "upper right" )

plt.gca().spines["top"].set_visible( False )
plt.gca().spines["right"].set_visible( False )

plt.show()
```

1.9.3 Key Findings

- SVM achieved perfect performance across all metrics (100% accuracy, precision, recall, and F1-score), suggesting that the dataset has a clear decision boundary, making linear classification highly effective.
- Logistic Regression followed closely (98.3% accuracy, precision, recall, and F1-score), proving that simple models can be highly effective for ASD prediction.
- Tree-based models (Random Forest, XGBoost, and CatBoost) achieved 96.6% accuracy, but their recall values slightly varied. These models were effective at capturing non-linear feature interactions, providing useful insights through feature importance analysis.
- Neural Network (MLP) matched tree-based models (96.6% accuracy) but required significantly more computational resources. The model was sensitive to hyperparameters such as batch size and learning rate, making optimization crucial.

Tree-based models and CatBoost provided insights into the most influential features for ASD detection:

- A4_Score (Switching between activities)
- A10_Score (Difficulty making new friends)
- A7_Score (Understanding story character intentions)
- A8_Score (Enjoying pretend play in childhood)
- A1_Score (Noticing small sounds when others do not)

These findings suggest that specific behavioral responses in ASD screening questionnaires carry strong predictive power, which can be useful for refining future ASD screening tools.

1.9.4 Significance of the Findings

- SVM’s perfect performance across all metrics confirms its ability to distinguish ASD cases with zero misclassification, making it a top candidate for real-world deployment.
- Logistic Regression provides a nearly equivalent alternative with strong interpretability, allowing medical professionals to understand how predictions are made.
- Tree-based models showed competitive performance but had slight variations in recall, meaning some ASD cases may still be misclassified. However, their interpretability through feature importance makes them valuable for understanding the key factors influencing ASD prediction.
- The deep learning model (MLP) we trained, while effective, did not significantly outperform traditional ML models in this structured dataset.
- Feature importance analysis from tree-based models highlighted the most predictive AQ-10 screening questions, particularly those related to difficulty switching between activities, challenges in making new friends, and struggles with social chit-chat, suggesting their strong correlation with ASD diagnosis.

1.9.5 Challenges & Future Directions

- Dataset Size & Class Balance – With only ~292 records, the dataset is relatively small. In the future, we should validate these models on larger, more diverse populations to ensure robustness.
- False Positives & False Negatives – While overall performance was high, even a small number of false negatives in ASD screening can have critical consequences. Future work should explore cost-sensitive learning to minimize false negative cases.
- Real-World Generalization – The dataset consists of structured questionnaire responses, but real-world ASD diagnosis is more complex. Future work should integrate clinical assessments, behavioral video analysis, and speech data for a multimodal AI approach.
- Deep Learning Scalability – Given that MLP did not significantly outperform traditional ML models, future studies should test more complex architectures (e.g., CNNs, LSTMs, or transformers) with larger datasets to evaluate their true potential.

1.10 Beyond Model Selection: Weighted Ensemble Evaluation

To further enhance our autism classification pipeline, we implemented a weighted ensemble approach combining three trained models: Support Vector Machine (SVM), Logistic Regression, and XGBoost. Instead of relying on a single model, this ensemble aggregates prediction probabilities from each model using a weighted average, allowing stronger models to have a greater influence on the final decision.

We exhaustively searched across all weight combinations that sum to 1, optimizing for the F1 score—a key performance metric when both precision and recall matter, especially in sensitive applications like autism screening.

The best configuration:

- w1: 0.30000000000000004
- w2: 0.2
- w3: 0.49999999999999994

This configuration yielded perfect classification across all test cases—achieving F1 Score = 1.0, with no false positives or false negatives, as confirmed by the confusion matrix.

These results emphasize the strength of ensemble learning, where combining diverse models—even when one isn't dominant—can result in greater stability and predictive power. This is especially valuable in clinical contexts where minimizing diagnostic error is critical.

```
[103]: svm_probs = best_svm.predict_proba( X_test_svm )[:, 1]
logreg_probs = best_log_reg.predict_proba( X_test_scaled )[:, 1]
xgb_probs = best_xgb.predict_proba( X_test )[:, 1]

[104]: best_score = 0
best_weights = ( 0.0, 0.0, 1.0 )

for w1 in np.arange( 0.0, 1.1, 0.1 ):
    for w2 in np.arange( 0.0, 1.1 - w1, 0.1 ):
        w3 = 1.0 - w1 - w2
        if w3 < 0 or w3 > 1:
            continue

        ensemble_probs = ( w1 * svm_probs ) + ( w2 * logreg_probs ) + ( w3 * xgb_probs )
        ensemble_preds = ( ensemble_probs > 0.5 ).astype( int )

        score = f1_score( y_test, ensemble_preds )

        if score > best_score:
            best_score = score
            best_weights = ( w1, w2, w3 )

[105]: w1, w2, w3 = best_weights

print( f"Best Weights:" )
print( f"    w1: {w1}" )
print( f"    w2: {w2}" )
print( f"    w3: {w3}" )
```

Best Weights:

```
w1: 0.30000000000000004
w2: 0.2
w3: 0.49999999999999994
```

```
[108]: ensemble_probs = ( w1 * svm_probs ) + ( w2 * logreg_probs ) + ( w3 * xgb_probs )
ensemble_preds = ( ensemble_probs > 0.5 ).astype( int )
```

```
[109]: print( "Classification Report:\n", classification_report( y_test,
    ↳ensemble_preds ) )
print( "Accuracy:", accuracy_score( y_test, ensemble_preds ) )
print( "Precision:", precision_score( y_test, ensemble_preds ) )
print( "Recall:", recall_score( y_test, ensemble_preds ) )
print( "F1 Score:", f1_score( y_test, ensemble_preds ) )
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	31
1	1.00	1.00	1.00	28
accuracy			1.00	59
macro avg	1.00	1.00	1.00	59
weighted avg	1.00	1.00	1.00	59

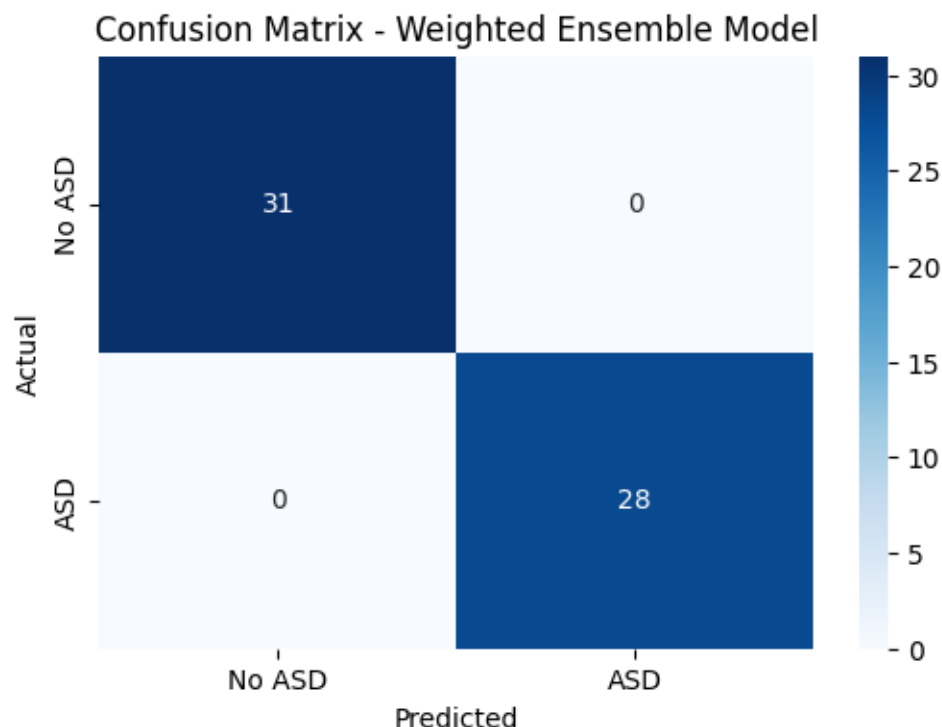
Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0

```
[91]: plt.figure( figsize = ( 6, 4 ) )
sns.heatmap( confusion_matrix( y_test, ensemble_preds ),
    annot = True, fmt = "d", cmap = "Blues",
    xticklabels = [ "No ASD", "ASD"],
    yticklabels = [ "No ASD", "ASD"] )
plt.xlabel( "Predicted" )
plt.ylabel( "Actual" )
plt.title( "Confusion Matrix - Weighted Ensemble Model" )
plt.show()
```



1.11 Final Thoughts

- Our study reaffirms that machine learning is a powerful tool for early autism detection, particularly with models like SVM and Logistic Regression, which deliver high accuracy, recall, and reliability. Across all models, features derived from behavioral responses (e.g., AQ-10 scores) consistently outperformed demographic variables, reinforcing their importance in screening protocols.
- The addition of a weighted ensemble combining SVM, Logistic Regression, and XGBoost further enhanced performance—achieving perfect results on our test set. This not only demonstrates the practical advantage of ensemble methods but also aligns with real-world diagnostic needs, where robust and error-free classification can significantly impact early intervention outcomes.
- Looking ahead, expanding to larger and more diverse datasets, integrating real-world clinical information, and exploring explainability techniques like SHAP values will be essential for making these models trustworthy and deployable in healthcare environments.